

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt

In [2]:
df = pd.read_excel('Linear regression project.xlsx')
df.head()
```

ID	Amount.Requested	Amount.Funded.By.Investors	Interest.Rate	Loan.Length	Loan.Purpose	Debt.To.Income.Ratio	Home.Ownership	Monthly.Inco	
0	81174.0	20000	20000.0	0.09	36 months	debt_consolidation	0.15	MORTGAGE	654
1	99592.0	19200	19200.0	0.12	36 months	debt_consolidation	0.28	MORTGAGE	458
2	80059.0	35000	35000.0	0.22	60 months	debt_consolidation	0.24	MORTGAGE	1150
3	15825.0	10000	9975.0	0.10	36 months	debt_consolidation	0.14	MORTGAGE	383
4	33182.0	12000	12000.0	0.12	36 months	credit_card	0.19	RENT	319

Out[2]:

In [3]: df.dtypes

Out[3]:

```
ID          float64
Amount.Requested      int64
Amount.Funded.By.Investors    float64
Interest.Rate        float64
Loan.Length         object
Loan.Purpose        object
Debt.To.Income.Ratio   float64
Home.Ownership       object
Monthly.Income       float64
Open.CREDIT.Lines     int64
Revolving.CREDIT.Balance int64
Inquiries.in.the.Last.6.Months int64
Employment.Length      object dtype:
object
```

In [4]: df.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2500 entries, 0 to 2499
Data columns (total 13 columns):
 #   Column           Non-Null Count Dtype  
 --- 
 0   ID               2499 non-null  float64 
 1   Amount.Requested  2500 non-null  int64  
 2   Amount.Funded.By.Investors  2500 non-null  float64 
 3   Interest.Rate    2500 non-null  float64 
 4   Loan.Length      2499 non-null  object  
 5   Loan.Purpose     2499 non-null  object  
 6   Debt.To.Income.Ratio  2500 non-null  float64 
 7   Home.Ownership   2499 non-null  object  
 8   Monthly.Income   2500 non-null  float64 
 9   Open.CREDIT.Lines 2500 non-null  int64  
 10  Revolving.CREDIT.Balance 2500 non-null  int64  
 11  Inquiries.in.the.Last.6.Months 2500 non-null  int64  
 12  Employment.Length      2422 non-null  object dtyp
es: float64(5), int64(4), object(4) memory usage: 254.0+KB
```

In [5]: df.isna().sum()

Out[5]:

```
ID               1
Amount.Requested  0
Amount.Funded.By.Investors  0
Interest.Rate     0
Loan.Length       1
Loan.Purpose      1
Debt.To.Income.Ratio  0
Home.Ownership    1
Monthly.Income    0
Open.CREDIT.Lines 0
Revolving.CREDIT.Balance 0
```

```
Inquiries.in.the.Last.6.Months 0
```

ID	Amount.Requested	Amount.Funded.By.Investors	Interest.Rate	Loan.Length	Loan.Purpose	Debt.To.Income.Ratio	Home.Ownership	Monthly
0	81174.0	20000	20000.00	0.09	36 months	debt_consolidation	0.15	MORTGAGE
1	99592.0	19200	19200.00	0.12	36 months	debt_consolidation	0.28	MORTGAGE
2	80059.0	35000	35000.00	0.22	60 months	debt_consolidation	0.24	MORTGAGE
Employment.Length	78	dtype:						1
int64								

```
In [6]:
```

```
df
```

```
Out[6]:
```

3	15825.0	10000	9975.00	0.10	36 months	debt_consolidation	0.14	MORTGAGE
4	33182.0	12000	12000.00	0.12	36 months	credit_card	0.19	RENT
...
2495	23735.0	30000	29950.00	0.17	60 months	debt_consolidation	0.19	MORTGAGE
2496	65882.0	16000	16000.00	0.14	60 months	home_improvement	0.22	OWN
2497	55610.0	10000	10000.00	0.14	36 months	debt_consolidation	0.05	MORTGAGE
2498	38576.0	6000	6000.00	0.12	36 months	major_purchase	0.17	RENT
2499	3116.0	9000	5242.75	0.14	36 months	debt_consolidation	0.07	RENT

2500 rows × 13 columns

In [7]:

```
df.nunique()
```

```
ID           2499
Amount.Requested      380
Amount.Funded.By.Investors    708
Interest.Rate        21
Loan.Length          3
Loan.Purpose         14
Debt.To.Income.Ratio 36
Home.Ownership       5
Monthly.Income        632
Open.CREDIT.Lines     29
Revolving.CREDIT.Balance 2349
Inquiries.in.the.Last.6.Months 10
Employment.Length      10 dtype:int64
```

In [8]:

```
gp=df.groupby('Amount.Requested').size()
gp.head()
```

```
Amount.Requested
1000   18
1125   1
1200   6
1400   3 1450   1 dtype:int64
```

In [9]:

```
df.drop_duplicates(subset=['Amount.Requested']
df
```

```
IDAmount.RequestedAmount.Funded.By.Investors
0 81174.0      20000
1 99592.0      19200
2 80059.0      35000
3 15825.0      10000
4 33182.0      12000
...
244779820.025875
```

,inplace=True)

		Interest.Rate	Loan.Length	Loan.Purpose	Debt.To.Income.Ratio	Home.Ownership
246692038.022750		0.09	36 months	debt_consolidation	0.15	MORTGAGE
247249533.017300		20000.0	0.12	36 months	debt_consolidation	MORTGAGE
248963256.0 19075 249178043.0 8475		19200.0	0.22	60 months	debt_consolidation	MORTGAGE
		35000.0	0.10	36 months	debt_consolidation	RENT
		9975.0 12000.0	0.12	36 months	credit_card	...
	
		25875.0	0.18	60 months	debt_consolidation	MORTGAGE
		22750.0	0.22	60 months	home_improvement	OWN
		17250.0	0.22	60 months	wedding	...
		19075.0	0.19	36 months	debt_consolidation	...
		8475.0	0.08	36 months	debt_consolidation	0.15
		0.16

In [10]:

```
df.describe()
```

Out[10]:

In [11]:

```
df.corr()
```

Out[11]:

In [12]:

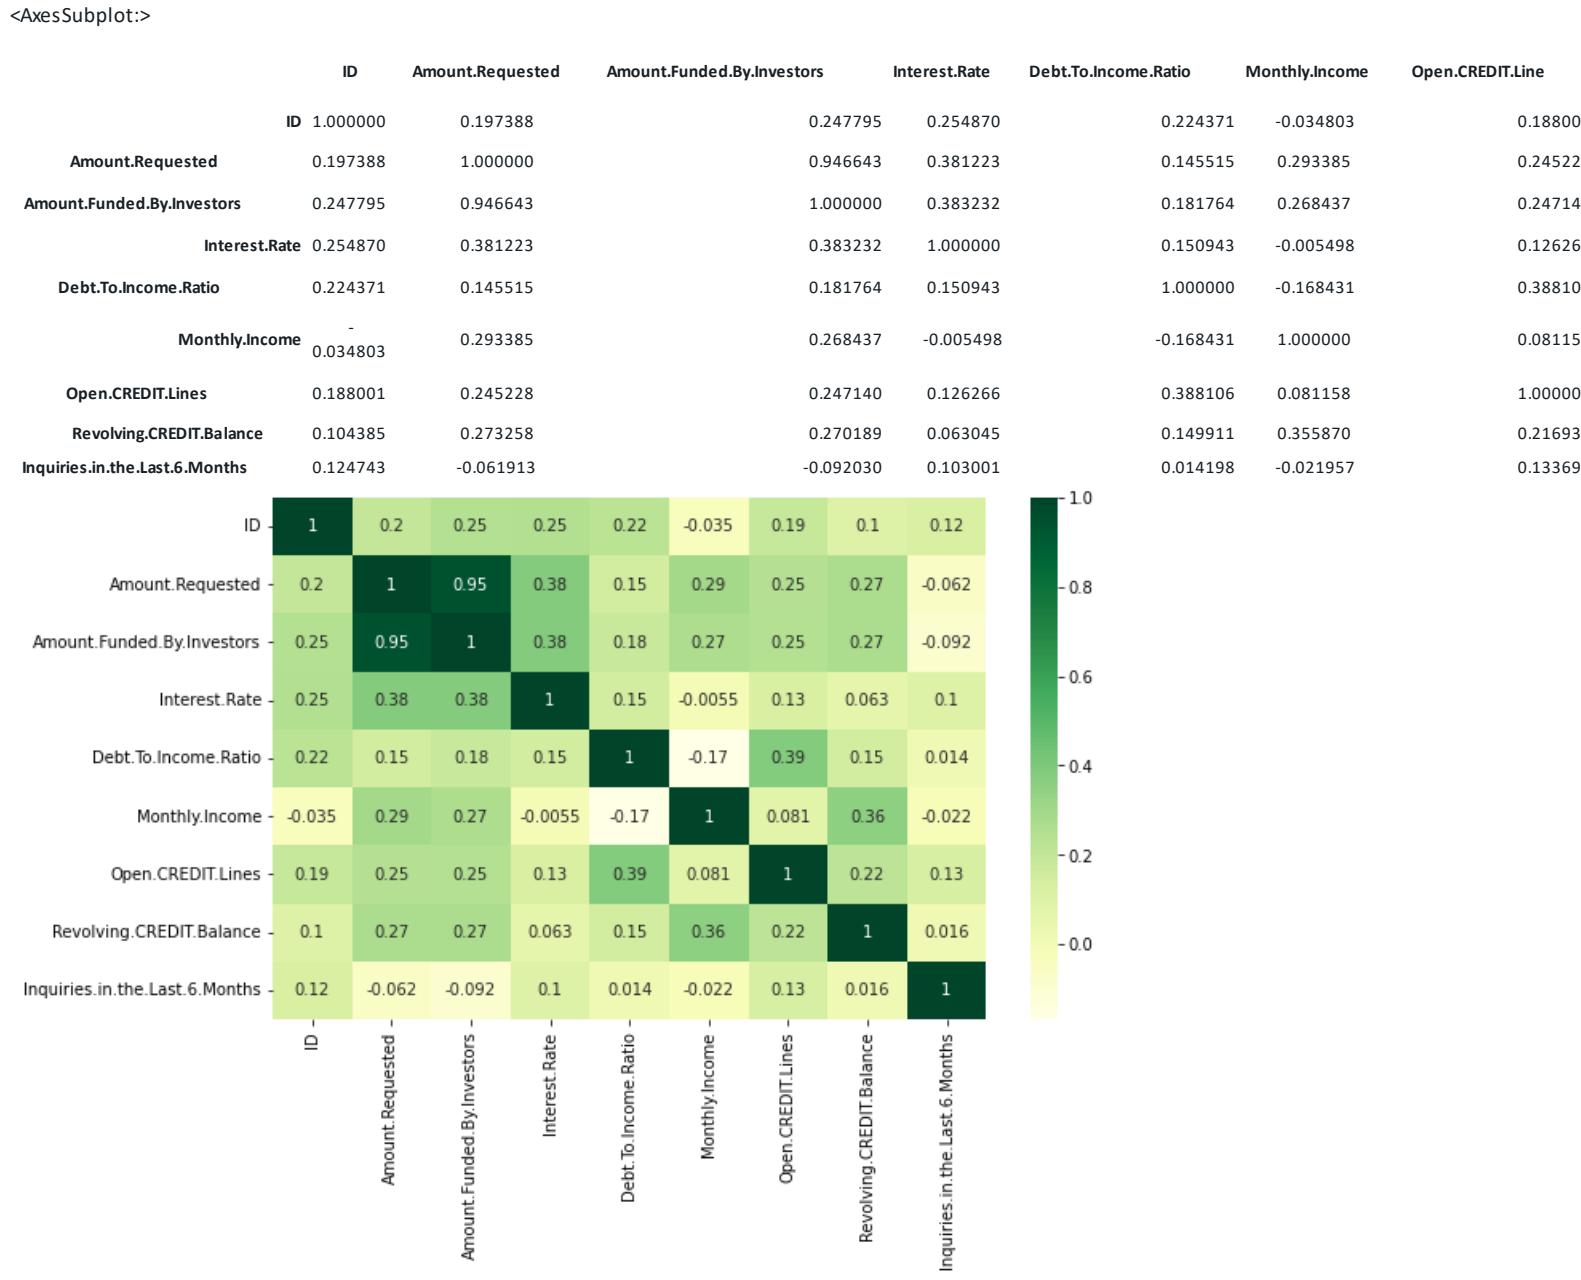
```
import seaborn as sns
```

In [13]:

	ID	Amount.Requested	Amount.Funded.By.Investors	Interest.Rate	Debt.To.Income.Ratio	Monthly.Income	Open.CREDIT.Lines	Revolving.CRE
count	380.000000	380.000000		380.000000	380.000000	380.000000	380.000000	380.000000
mean	56423.807895	13231.973684		12799.127105	0.136421	0.165658	5348.496895	10.057895
std	28890.272016	8239.491632		8310.901170	0.043292	0.075625	5682.379787	4.329053
min	17.000000	1000.000000		200.000000	0.050000	0.000000	666.670000	2.000000
25%	33037.750000	6562.500000		6056.250000	0.100000	0.110000	3333.330000	7.000000
50%	58928.000000	11912.500000		11362.500000	0.140000	0.160000	4583.330000	9.000000
75%	79733.000000	18631.250000		18406.250000	0.170000	0.220000	6025.000000	13.000000
max	103218.000000	35000.000000		35000.000000	0.240000	0.340000	102750.000000	25.000000

```
plt.figure(figsize=(10,6))
sns.heatmap(df.corr(),annot=True,cmap='YIGn')
```

Out[13]:

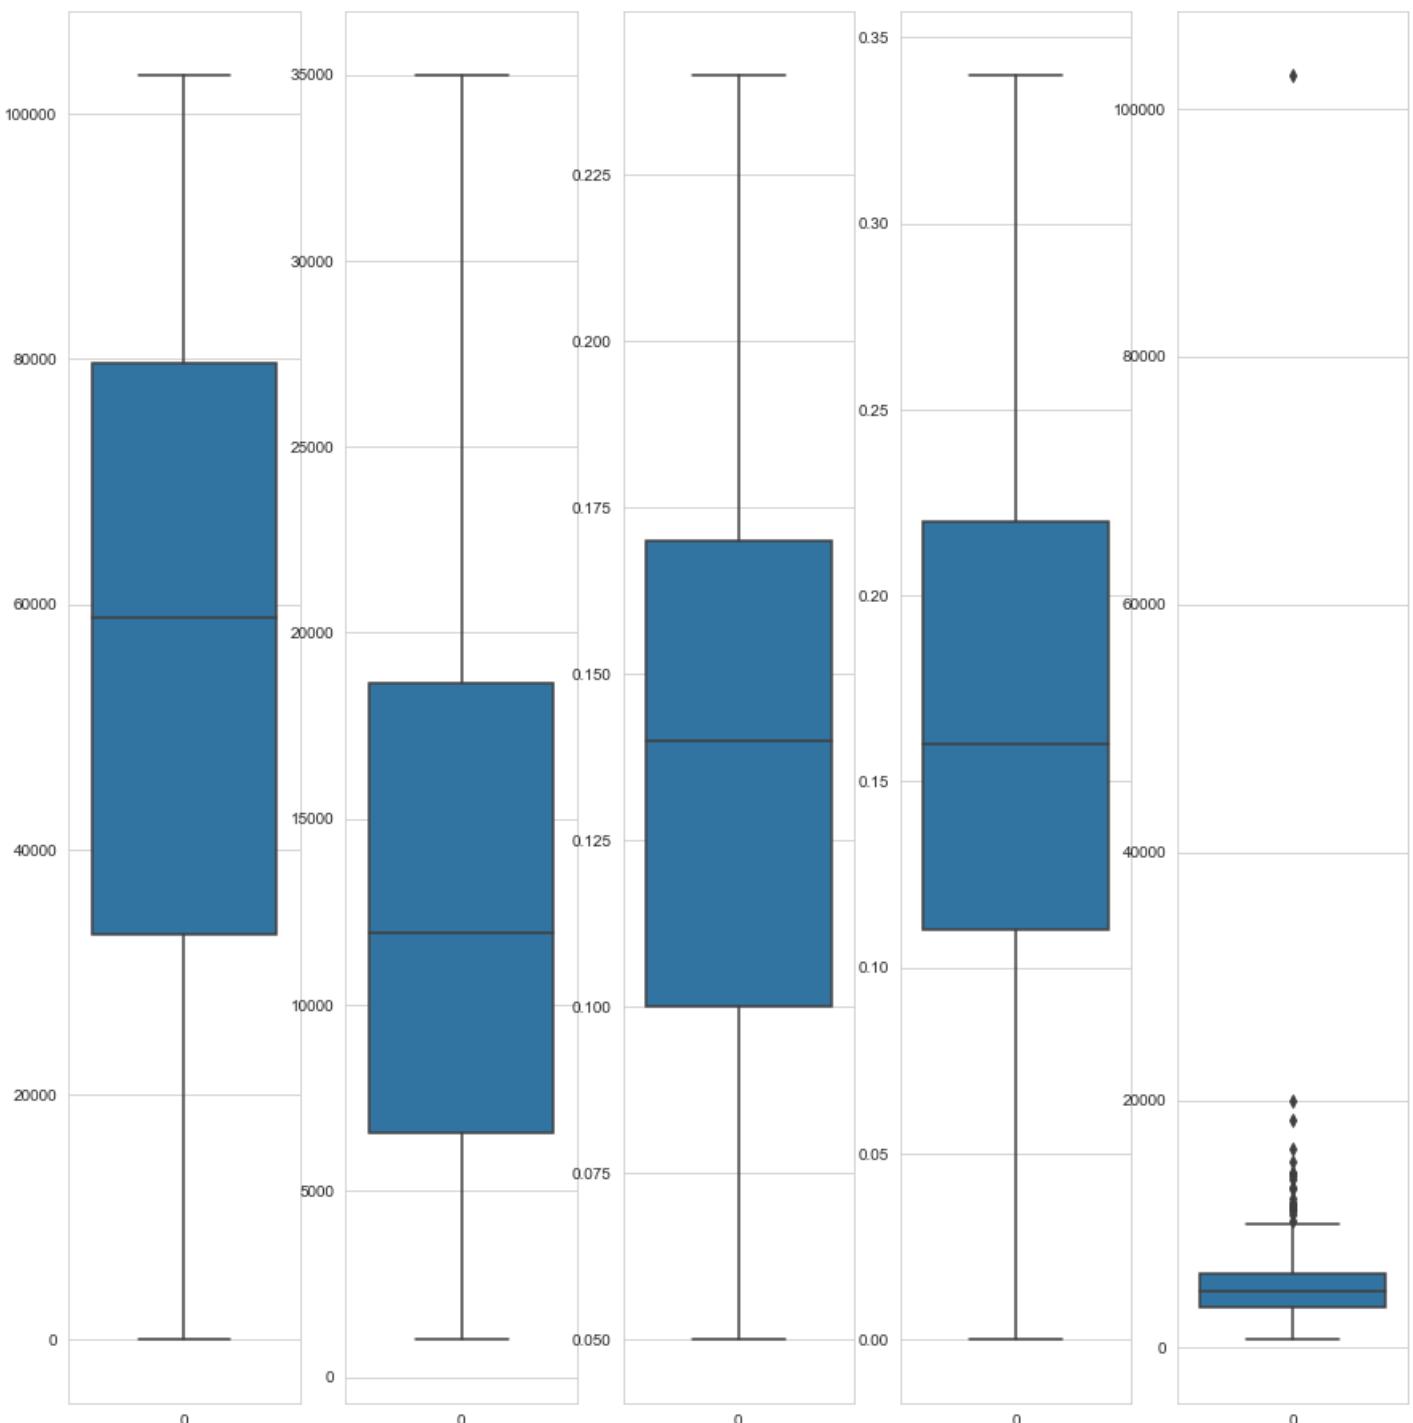


In [14]:

```
plt.figure(figsize=(15,16))
sns.set_style('whitegrid')
plt.subplot(1,5,1)
sns.boxplot(data=df['ID'])
plt.subplot(1,5,2)
sns.boxplot(data=df['Amount.Requested'])
plt.subplot(1,5,3)
sns.boxplot(data=df['Interest.Rate'])
plt.subplot(1,5,4)
sns.boxplot(data=df['Debt.To.Income.Ratio'])
plt.subplot(1,5,5)
sns.boxplot(data=df['Monthly.Income'])
```

Out[14]:

<AxesSubplot:>



In [15]:

`df.shape`

(380 13)

5)

Out[15]:

```
2489 False False False False False False False False
```

380 rows × 13 columns



In [17]:

```
x=pd.DataFrame(data=df['Amount.Requested'],columns=('Amount.Requested')) x.head()
```

Out[17]: Amount.Requested

```
0    20000  
1    19200  
2    35000  
3    10000  
4    12000
```

In [18]:

```
y=pd.DataFrame(data=df['Interest.Rate'],columns=('Interest.Rate')) y.head()
```

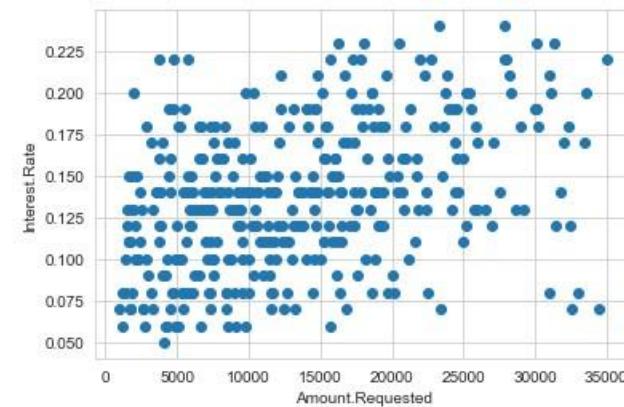
Out[18]:

```
Interest.Rate  
0    0.09  
1    0.12  
2    0.22  
3    0.10  
4    0.12
```

In [19]:

```
plt.scatter(x,y) plt.xlabel("Amount.Requested")  
plt.ylabel("Interest.Rate")
```

Out[19]: Text(0, 0.5, 'Interest.Rate')



In [20]:

```
from sklearn.model_selection import train_test_split x_train,x_test,y_train,y_test =  
train_test_split(x,y,train_size=0.2,  
random_state=4)
```

In [21]:

```
from sklearn.linear_model import LinearRegression
```

In [22]:

```
lr = LinearRegression()
```

In [23]:

```
lr.fit(x_train,y_train)
```

Out[23]:

```
LinearRegression  
LinearRegression()
```

In [24]:

```
x_test
```

Out[24]: Amount.Requested

```
2472          17300
17           14000
1186    7900 29   14500
316     31075 ... ...
393           21250
10            2000
170            11500
1508            31500
42             5500
```

304 rows × 1 columns

In [25]:

```
y_test
```

Out[25]: `Interest.Rate`

```
2472      0.22      17
           0.10
1186      0.14
29        0.08
316      0.20      ...
           ...
           393
           0.19
10        0.20
170       0.11
1508       0.12
42        0.19
```

304 rows × 1 columns

In [26]:

```
x_train
```

Out[26]: `Amount.Requested`

```
2380          15575
163            2800
1629    14100 538  1500
80     7450 ... ...
218            15200
2239            5175
743            19600
604    2400 382  18400
```

76 rows × 1 columns

In [27]:

```
y_train
```

Out[27]: `Interest.Rate`

```
2380      0.14
163       0.10
1629      0.19
538       0.13
```

```
80      0.13
...
...     218
0.11
2239      0.18
743       0.21
604       0.11
382       0.19
```

76 rows × 1 columns

In [28]:

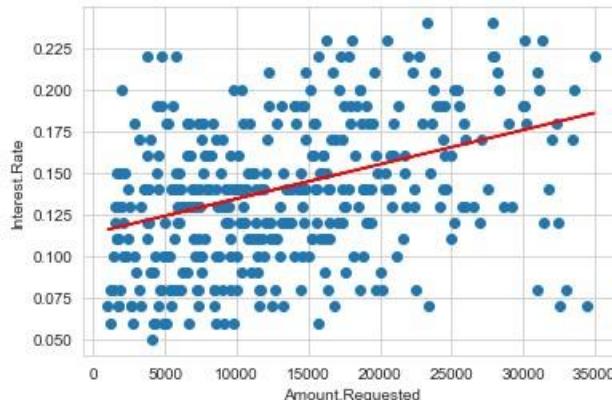
```
lr.predict(x_test.iloc[0].values.reshape(1,1))
```

```
C:\Users\kiransindam\AppData\Local\Programs\Python\Python310\lib\site-packages\sklearn\base.py:450: UserWarning: X does not have valid feature names
, but LinearRegression was fitted with feature names  warnings.warn(
```

```
array([[0.14950488]])
```

In [29]:

```
plt.scatter(df['Amount.Requested'],df['Interest.Rate'])
plt.plot(x_test,lr.predict(x_test),color='red')
plt.xlabel('Amount.Requested') plt.ylabel('Interest.Rate')
```



Out[29]: Text(0, 0.5, 'Interest.Rate')

In [30]:

```
lr.coef_
```

```
array([[2.0621461e-06]])
```

In [31]:

```
lr.intercept_
```

```
array([0.11382976])
```

In [32]:

```
#y=mx + b
#m*17300 +b
```

In [33]:

```
y_predicte=lr.predict(x_train)y_predicte
```

Out[30]:

Out[31]:

```
array([[0.14594768],  
[0.11960377],  
[0.14290602],  
[0.11692298],  
[0.12919274],  
[0.16409457],  
[0.12955362],  
[0.12187213],  
[0.12310941],  
[0.13393568],  
[0.1334717],  
[0.18054018],  
[0.13919415],  
[0.14331845],  
[0.12816167],  
[0.12166591],  
[0.14599924],  
[0.15862988],  
[0.17120897],  
[0.14068921],  
[0.14744274],  
[0.14780361],  
[0.15383539],  
[0.13032693],  
[0.11769628],  
[0.11841803],  
[0.12898653],  
[0.11774783],  
[0.13909105],  
[0.13934881],  
[0.12847099],  
[0.15754725],  
[0.13527608],  
[0.12769769],  
[0.1525981],  
[0.13651336],  
[0.12939896],  
[0.14764895],  
[0.13445122],  
[0.1461539],  
[0.11718074],  
[0.15589754],  
[0.14733963],  
[0.15223723],  
[0.12785235],  
[0.12321252],  
[0.14888624],  
[0.12403738],  
[0.14007057],  
[0.15316519],  
[0.13775065],  
[0.16265107],  
[0.11836648],  
[0.14367932],  
[0.12867721],  
[0.16229019],  
[0.15507268],  
[0.11712919],  
[0.13455432],  
[0.14228737],  
[0.14476195],  
[0.14352466],  
[0.16074358],  
[0.13125489],  
[0.15074217],  
[0.14115319],  
[0.14986576],  
[0.1215628],  
[0.16703313],  
[0.13017226],  
[0.17590035],  
[0.14517438],  
[0.12450136],
```

```
[0.15424782],  
[0.11877891],  
[0.15177324]])
```

In [34]:

```
from sklearn.metrics import mean_squared_error,r2_score
```

In [35]:

```
mse_train = mean_squared_error(y_train,y_predicte)  
mse_train
```

Out[35]: 0.001417332913089926

In [36]:

```
score_train = r2_score(y_train,y_predicte) score_train
```

Out[36]:

0.14031578165778436

In [37]:

```
y_predicted = lr.predict(x_test)y_predicted
```

Out[37]:

```
array([[0.14950488],  
       [0.1426998 ],  
       [0.13012071],  
       [0.14373087],  
       [0.17791095],  
       [0.12888342],  
       [0.16435234],  
       [0.13439966],  
       [0.16332126],  
       [0.14574147],  
       [0.15301053],  
       [0.15713482],  
       [0.13960658],  
       [0.12965673],  
       [0.17156985],  
       [0.13151266],  
       [0.14440107],  
       [0.12908964],  
       [0.12661506],  
       [0.13352325],  
       [0.12352184],  
       [0.16579584],  
       [0.11785094],  
       [0.12048018],  
       [0.16125912],  
       [0.12372806],  
       [0.14161718],  
       [0.13177043],  
       [0.1257902 ],  
       [0.12584176],  
       [0.15254655],  
       [0.1646101 ],  
       [0.12594486],  
       [0.13460588],  
       [0.13383257],  
       [0.12599642],  
       [0.15058751],  
       [0.12151125],  
       [0.14388554],  
       [0.12414049],  
       [0.13898794],  
       [0.17837493],  
       [0.1243467 ],  
       [0.18291165],  
       [0.14847381],  
       [0.13130644],  
       [0.13084246],  
       [0.14527748],  
       [0.15094839],  
       [0.16187776],  
       [0.15141237],  
       [0.13140955],  
       [0.14697876],  
       [0.14285446],  
       [0.12491379],  
       [0.16538341],  
       [0.1256871 ],  
       [0.11929444],
```

[0.1532683],
[0.17069344],
[0.1378022],
[0.13610093],
[0.12609953],
[0.15197946],
[0.12285165],
[0.14079232],
[0.12723371],
[0.13831774],
[0.14636011],
[0.13795687],
[0.14166873],
[0.15342296],
[0.13738978],
[0.14754585],
[0.14393709],
[0.15264966],
[0.1567224],
[0.15053596],
[0.13522452],
[0.14419486],
[0.13362636],
[0.12914119],
[0.14558681],
[0.14218427],
[0.14785517],
[0.1321313],
[0.12393427],
[0.13274995],
[0.14105009],
[0.11743851],
[0.12475913],
[0.15837211],
[0.1355854],
[0.15857833],
[0.12522311],
[0.164507],
[0.13238907],
[0.12130504],
[0.14187494],
[0.17399287],
[0.12362495],
[0.15455714],
[0.16022804],
[0.17945756],
[0.15089683],
[0.15017508],
[0.12826478],
[0.15136082],
[0.14063766],
[0.14929867],
[0.16636293],
[0.11754162],
[0.13259529],
[0.14027678],
[0.14723652],
[0.17363199],
[0.11749007],
[0.1376991],
[0.15888765],
[0.13692579],
[0.1278008],
[0.13305927],
[0.1638368],
[0.17198228],
[0.16718779],
[0.13878172],
[0.12836789],
[0.15919697],
[0.1441433],
[0.12486224],
[0.12063484],
[0.14795827],
[0.15584598],
[0.12104727],

[0.14641166],
[0.12805856],
[0.12893498],
[0.12553244],
[0.17569414],
[0.11939755],
[0.12125348],
[0.12207834],
[0.13599783],
[0.14022523],
[0.15228878],
[0.13805997],
[0.16368214],
[0.14383398],
[0.13486365],
[0.11846959],
[0.13156421], [0.16522875],
[0.1849738],
[0.12728526],
[0.16280573],
[0.12749147],
[0.1758488],
[0.1214597],
[0.15646463],
[0.12429515],
[0.13723511],
[0.11913978],
[0.1299145],
[0.12635729],
[0.1186758],
[0.15404161],
[0.1285741],
[0.13718356],
[0.16208398],
[0.1497111],
[0.17775629],
[0.13517297],
[0.12929585],
[0.13671958],
[0.14295757],
[0.13506986],
[0.15703172],
[0.14001901],
[0.167394],
[0.15651618],
[0.12171747],
[0.13852396],
[0.11975843],
[0.1158919],
[0.1328015],
[0.16002183],
[0.11728385],
[0.13620404],
[0.12280009],
[0.12053173],
[0.12001619],
[0.18311787],
[0.13233752],
[0.12671817],
[0.12981139],
[0.13342014],
[0.15450559],
[0.16847663],
[0.14703031],
[0.12506845],
[0.13053314],
[0.13331704],
[0.131358],
[0.12331563],
[0.12249077],
[0.14424641],
[0.13537918],
[0.14682409],
[0.16414612],
[0.12548088],
[0.13434811],

[0.13857551],
[0.12274854],
[0.14991731],
[0.17141519],
[0.1695077],
[0.12702749],
[0.1229032],
[0.12743992],
[0.1384724],
[0.15971251],
[0.12326407],
[0.14538059],
[0.16744556],
[0.14512282],
[0.17615812],
[0.16631137],
[0.1250169],
[0.15533045],
[0.1722916],
[0.12496535],
[0.15466025],
[0.12764614], [0.18600487],
[0.14084387],
[0.13197664],
[0.13290461],
[0.12042862],
[0.12254232],
[0.17785939],
[0.11681987],
[0.12620263],
[0.15991872],
[0.17981843],
[0.16301194],
[0.14610234],
[0.11702608],
[0.14878313],
[0.14579302],
[0.14955644],
[0.13548229],
[0.1454837],
[0.12537777],
[0.12945051],
[0.15012353],
[0.14156562],
[0.12738837],
[0.13115178],
[0.15610375],
[0.11898512],
[0.13574006],
[0.14656633],
[0.14409175],
[0.1483707],
[0.12573865],
[0.13326548],
[0.12754303],
[0.11810871],
[0.11614967],
[0.12377961],
[0.14110164],
[0.152495],
[0.11630433],
[0.18064329],
[0.15388694],
[0.122233],
[0.18095261],
[0.17285869],
[0.12455292],
[0.12212989],
[0.13099712],
[0.1173354],
[0.14960799],
[0.16961081],
[0.13661647],
[0.13403879],
[0.11816026],
[0.12883187],

```
[0.11671676],  
[0.12852255],  
[0.14429797],  
[0.14775206],  
[0.13733822],  
[0.16584739],  
[0.13249218],  
[0.15311364],  
[0.18188058],  
[0.12640885],  
[0.15208257],  
[0.15765036],  
[0.11795405],  
[0.13754444],  
[0.17878736],  
[0.12517156]))
```

In [38]:

```
mse_test=mean_squared_error(y_train,y_predicte)  
mse_test
```

Out[38]: 0.001417332913089926

In [39]:

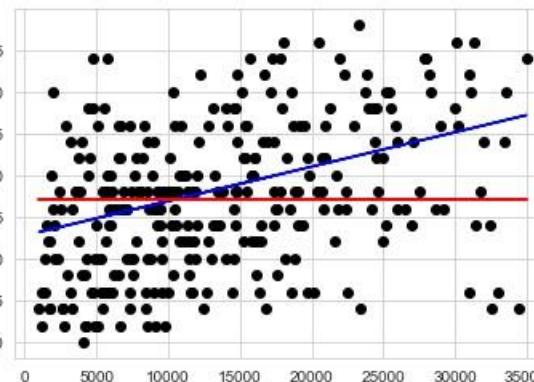
```
score_test=r2_score(y_train,y_predicte) score_test
```

Out[39]:

0.14031578165778436

In [40]:

```
plt.scatter(x_test, y_test, color='black')  
plt.plot(x_test,y_predicted , color='blue') a  
= y_test.to_numpy() k = np.mean(a)  
plt.plot(x_test,[k]*len(a), color='red') plt.show()
```



In [41]:

```
import statsmodels.api as sm x2 =  
sm.add_constant(x)  
est=sm.OLS(y,x2)  
est2=est.fit() est2.summary()
```

Out[41]: OLS Regression Results

Dep. Variable: Interest.Rate R-squared: 0.145

Model: OLS Adj. R-squared: 0.143

Method: Least Squares F-statistic: 64.28

Date: Sun, 04 Sep Prob (F-statistic): 1.36e-14
2022

Time: 18:45:08 Log-Likelihood: 684.26

No. Observations: 380 AIC: -1365. Df Residuals: 378

BIC: -1357.

Df Model: 1

Covariance Type: nonrobust

	coef	std err	t	P> t	[0.025	0.975]
--	------	---------	---	------	--------	--------

const	0.1099	0.004	28.234	0.000	0.102	0.118
-------	--------	-------	--------	-------	-------	-------

Amount.Requested	2.003e-06	2.5e-07	8.017	0.000	1.51e-06	2.49e-06
------------------	-----------	---------	-------	-------	----------	----------

Omnibus:	4.723	Durbin-Watson:	1.992
Mnibus):	0.094	Jarque-Bera	3.337
		(JB):	
Skew:	0.066	Prob(JB):	0.189
Kurtosis:	2.560	Cond. No.	2.95e+04

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 2.95e+04. This might indicate that there are strong multicollinearity or other numerical problems.

In []:

In []: